

# Operating Systems I

Unit 4 – Processes

Prof. Dr. Alejandro Zunino



# Module 4: Processes

- Process Concept
- Process Scheduling
- Operation on Processes
- Cooperating Processes
- Interprocess Communication



# **Process Concept**

- An operating system executes a variety of programs:
  - Batch system jobs
  - Time-shared systems user programs or tasks
- The terms *job* and *process* are used almost interchangeably.
- Process: "a program in execution"
  - process execution must progress in sequential fashion.
- A process includes:
  - program counter
  - stack
  - data section



# Major Requirements of an Operating System

- Interleave the execution of several processes to maximize processor utilization while providing reasonable response time
- Allocate resources to processes
- Support interprocess communication and user creation of processes



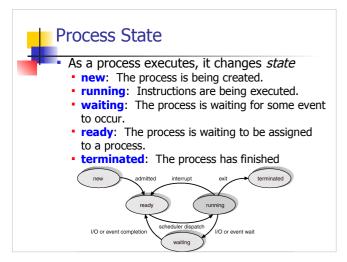
#### **Process Creation**

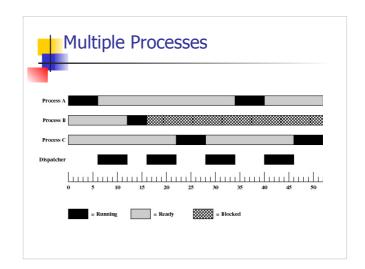
- Submission of a batch job
- User clicks on an application or file
- Create to provide a service such as printing
- Spawned by an existing process



#### **Process Termination**

- Batch job issues Halt instruction
- User logs off
- Process executes a service request to terminate
- Parent terminates so child processes terminate
- Operating system intervention
  - E.g. when deadlock occurs
- Error and fault conditions
  - E.g. memory unavailable, protection error, arithmetic error, I/O failure, invalid instruction





# Scheduler

- Program that moves the processor from one process to another
- Prevents a single process from monopolizing processor time.
  - Every process should be able to use the processor for a fair amount of time

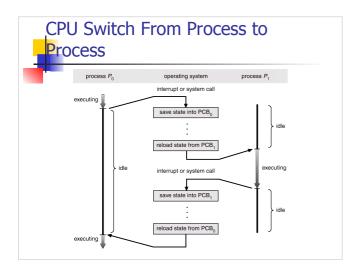


# Process Control Block (PCB)

Information associated with each process

- Process state
- Program counter
- CPU registers
- CPU scheduling information
- Memory-management inf.
- Accounting information
- I/O status information

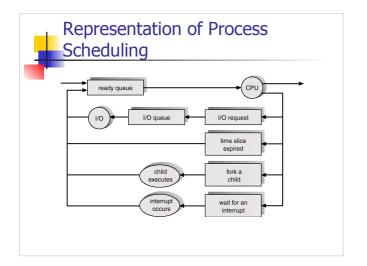
ch process.	
pointer	process state
process number	
program counter	
registers	
memory limits	
list of open files	
	:
	pointer process program regi

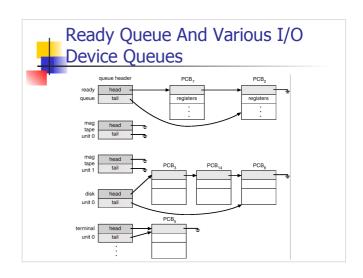


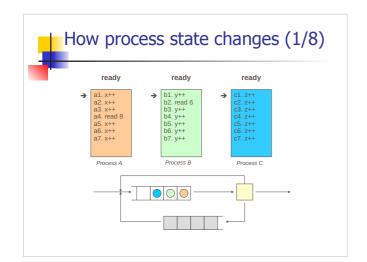


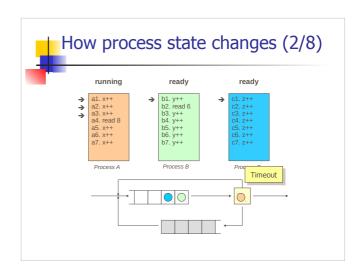
# Process Scheduling Queues

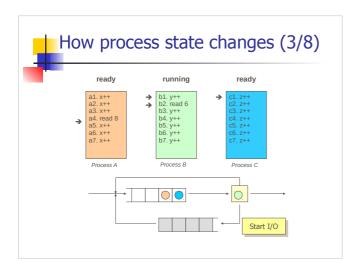
- Job queue set of all processes in the system.
- Ready queue set of all processes residing in main memory, ready and waiting to execute.
- Device queues set of processes waiting for an I/O device.
- Process migration between the various queues.

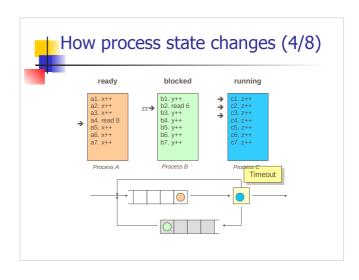


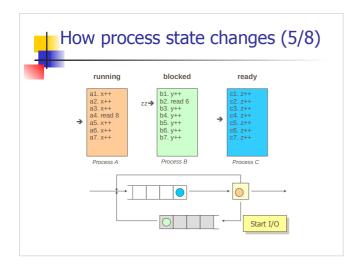


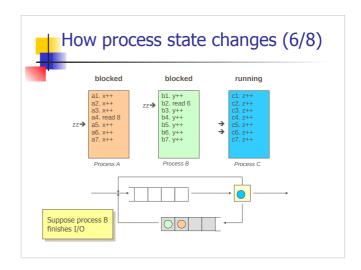


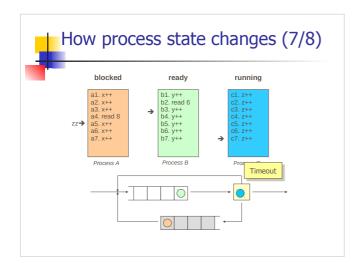


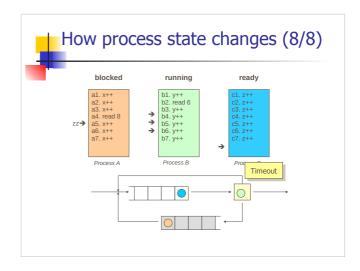


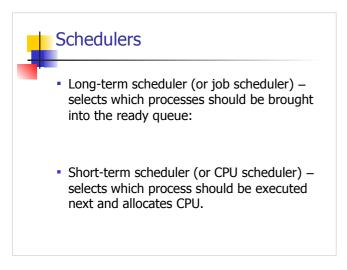


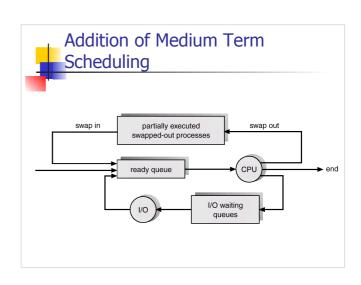














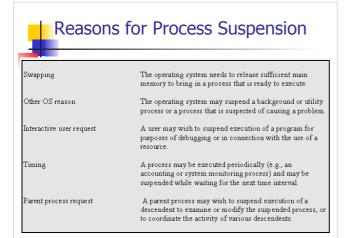
# Schedulers (Cont.)

- Short-term scheduler is invoked very frequently (milliseconds)
   b (must be fast).
- Long-term scheduler is invoked very infrequently (seconds, minutes) P (may be slow).
- The long-term scheduler controls the degree of multiprogramming.
- Processes can be described as either:
  - I/O-bound process spends more time doing I/O than computations, many short CPU bursts.
  - CPU-bound process spends more time doing computations; few very long CPU bursts.



### **Context Switch**

- When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process.
- Context-switch time is overhead; the system does no useful work while switching.
- Time dependent on hardware support.





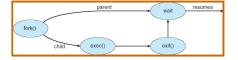
#### **Process Creation**

- Parent process creates children processes, which, in turn create other processes, forming a tree of processes.
- Resource sharing
  - Parent and children share all resources.
  - Children share subset of parent's resources.
  - Parent and child share no resources.
- Execution
  - Parent and children execute concurrently.
  - Parent waits until children terminate.



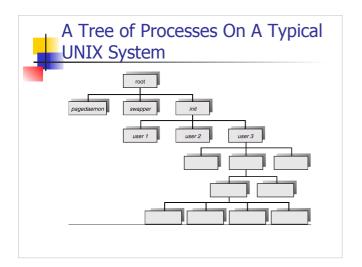
## Process Creation (Cont.)

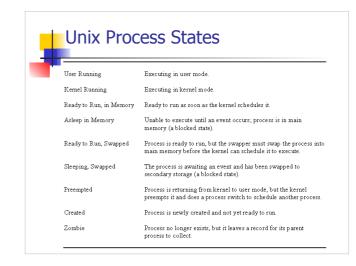
- Address space
  - Child duplicate of parent.
  - Child has a program loaded into it.
- UNIX examples
  - fork system call creates new process
  - execve system call used after a fork to replace the process' memory space with a new program.

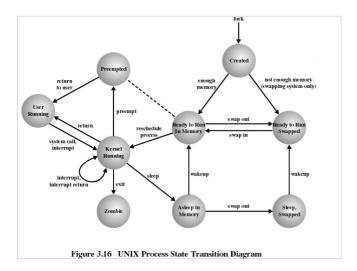




#### **Process Creation Example**









#### **Process Termination**

- Normal completion
- Time limit exceeded
- Memory unavailable
- Bounds violation
- Protection error
  - example write to read-only file
- Arithmetic error
- Time overrun
  - process waited longer than a specified maximum for an event



#### **Process Termination**

- I/O failure
- Invalid instruction
  - happens when try to execute data
- Privileged instruction
- Data misuse
- · Operating system intervention
  - such as when deadlock occurs
- Parent terminates so child processes terminate
- Parent request



# **IPC & Cooperative Processes**



# **Cooperating Processes**

- Independent process cannot affect or be affected by the execution of another process.
- Cooperating process can affect or be affected by the execution of another process
- Advantages of process cooperation
  - Information sharing
  - Computation speed-up
  - Modularity
  - Convenience



### **Producer-Consumer Problem**

- Paradigm for cooperating processes, producer process produces information that is consumed by a consumer process.
  - unbounded-buffer places no practical limit on the size of the buffer.
  - bounded-buffer assumes that there is a fixed buffer size



# Bounded-Buffer – Shared-Memory Solution

Shared data

var n; type item = ...; var buffer. array [0..n-1] of item; in, out: 0..n-1;

Producer process
 repeat

produce an item in *nextp* ... while *in*+1 mod *n* = out do *no-op*;

buffer [in] := nextp;

in :=in+1 mod η;
until false;



# Bounded-Buffer (Cont.)

Consumer process

#### repeat

while in = out do no-op; nextc := buffer [out]; out := out+1 mod n;

consume the item in nexto

until false

 Solution is correct, but can only fill up n-1 buffer.



# Interprocess Communication (IPC)

- Mechanism for processes to communicate and to synchronize their actions.
- Message system processes communicate with each other without resorting to shared variables.
- IPC facility provides two operations:
  - **send**(*message*) message size fixed or variable
  - receive(message)



# Interprocess Communication (IPC)

- If P and Q wish to communicate, they need to:
  - establish a communication link between them
  - exchange messages via send/receive
- Implementation of communication link
  - physical (e.g., shared memory, hardware bus)
  - logical (e.g., logical properties)



# **Implementation Questions**

- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?



### **Direct Communication**

- Processes must name each other explicitly:
  - send (P, message) send a message to process P
  - receive(Q, message) receive a message from process Q
- Properties of communication link
  - Links are established automatically.
  - A link is associated with exactly one pair of communicating processes.
  - Between each pair there exists exactly one link.
  - The link may be unidirectional, but is usually bidirectional.



#### **Indirect Communication**

- Messages are directed and received from mailboxes (also referred to as ports).
  - Each mailbox has a unique id.
  - Processes can communicate only if they share a mailbox.
- Properties of communication link
  - Link established only if processes share a common mailbox
  - A link may be associated with many processes.
  - Each pair of processes may share several communication links.
  - Link may be unidirectional or bi-directional.



#### **Indirect Communication**

- Operations
  - · create a new mailbox
  - send and receive messages through mailbox
  - destroy a mailbox



# Indirect Communication (Continued)

- Mailbox sharing
  - $P_1$ ,  $P_2$ , and  $P_3$  share mailbox A.
  - P<sub>1</sub>, sends; P<sub>2</sub> and P<sub>3</sub> receive.
  - Who gets the message?
- Solutions
  - Allow a link to be associated with at most two processes.
  - Allow only one process at a time to execute a receive operation.
  - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.



## **Synchronization**

- Message passing may be either blocking or non-blocking.
- Blocking is considered synchronous
- Non-blocking is considered asynchronous
- send and receive primitives may be either blocking or non-blocking.



# Buffering

- Queue of messages attached to the link; implemented in one of three ways.
  - 1. Zero capacity 0 messages Sender must wait for receiver (rendezvous).
  - 2. Bounded capacity finite length of *n* messages Sender must wait if link full.
  - 3. Unbounded capacity infinite length Sender never waits.



# Exception Conditions – Error Recovery

- Process terminates
- Lost messages
- Scrambled Messages



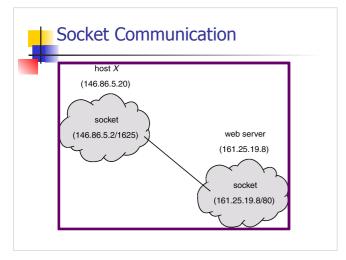
# **Client/Server Communication**

- Sockets
- Remote Procedure Calls
- Remote Method Invocation (Java)



#### **Sockets**

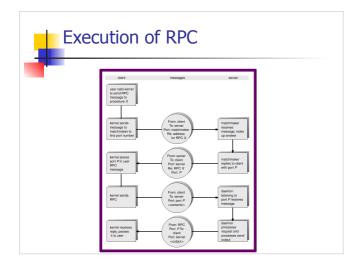
- A socket is defined as an endpoint for communication.
- Concatenation of IP address and port
- The socket 161.25.19.8:1625 refers to port 1625 on host 161.25.19.8
- Communication consists between a pair of sockets.





#### Remote Procedure Calls

- Remote procedure call (RPC) abstracts procedure calls between processes on networked systems.
- Stubs client-side proxy for the actual procedure on the server.
- The client-side stub locates the server and *marshalls* the parameters.
- The server-side stub receives this message, unpacks the marshalled parameters, and peforms the procedure on the server.





# **Remote Method Invocation**

- Remote Method Invocation (RMI) is a Java mechanism similar to RPCs.
- RMI allows a Java program on one machine to invoke a method on a remote object.

